

Ex.No – 7

Roll No – 210701149

Implement Linear and Logistic Regression

AIM:

To implement linear and logistic regression techniques in machine learning.

PROCEDURES:

Linear Regression

1. Define vectors for heights and weights.
2. Combine the heights and weights into a data frame.
3. Fit a linear regression model using height to predict weight.
4. Print the summary of the linear regression model to view model statistics.
5. Open a new graphical device for plotting.
6. Create a scatter plot of height vs. weight data points.
7. Label the plot with a title, x-axis label (Height), and y-axis label (Weight).
8. Set plot points with specific color (blue) and style (solid circle).
9. Add the fitted linear regression line to the plot.
10. Customize the regression line with red color and a thicker width.

Logistic Regression

1. Load the `mtcars` dataset.
2. Convert the `am` column from numeric to a factor with labels "Automatic" and "Manual."

3. Fit a logistic regression model to predict `am` (transmission) based on `mpg` (miles per gallon).
4. Print the summary of the logistic regression model.
5. Predict the probabilities of manual transmission using the logistic model.
6. Print the predicted probabilities for manual transmission.
7. Create a scatter plot of `mpg` vs. transmission type (manual/automatic).
8. Label the plot with a title, x-axis label (MPG), and y-axis label (Probability of Manual Transmission).
9. Set plot points with blue color and solid circles.
10. Add the logistic regression curve to the plot, colored red with a thicker line.

CODE:

```
LinearRegression.py
```

```
# Sample data
```

```
heights <- c(150, 160, 165, 170, 175, 180, 185)
```

```
weights <- c(55, 60, 62, 68, 70, 75, 80)
```

```
# Create a data frame
```

```
data <- data.frame(heights, weights)
```

```
# Fit a linear regression model
```

```
linear_model <- lm(weights ~ heights, data = data)
```

```
# Print the summary of the model
```

```
print(summary(linear_model))
```

```
# Plotting the data and regression line
```

```
dev.new()
```

```
plot(data$heights, data$weights,
```

```
      main = "Linear Regression: Weight vs. Height",
```

```
      xlab = "Height (cm)",
```

```
      ylab = "Weight (kg)",
```

```
      pch = 19, col = "blue")
```

```
# Add regression line
```

```
abline(linear_model, col = "red", lwd = 2)
```

```
LogisticRegression.py
```

```
# Load the dataset
```

```
data(mtcars)
```

```
# Convert 'am' to a factor (categorical variable)
```

```
mtcars$am <- factor(mtcars$am, levels = c(0, 1),
```

```
                  labels = c("Automatic", "Manual"))
```

```
# Fit a logistic regression model
```

```
logistic_model <- glm(am ~ mpg, data = mtcars, family = binomial)
```

```
# Print the summary of the model
```

```
print(summary(logistic_model))
```

```

# Predict probabilities for the logistic model
predicted_probs <- predict(logistic_model, type = "response")
# Display the predicted probabilities
print(predicted_probs)
# Plotting the data and logistic regression curve
plot(mtcars$mpg, as.numeric(mtcars$am) - 1,
     main = "Logistic Regression: Transmission vs. MPG",
     xlab = "Miles Per Gallon (mpg)",
     ylab = "Probability of Manual Transmission",
     pch = 19, col = "blue")
# Add the logistic regression curve
curve(predict(logistic_model, data.frame(mpg = x), type = "response"),
      add = TRUE, col = "red", lwd = 2)

```

OUTPUT:

```

Call:
lm(formula = weights ~ heights, data = data)


Residuals:
    1      2      3      4      5      6      7 
1.7049 -0.4754 -2.0656  0.3443 -1.2459  0.1639  1.5738 

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -54.40984     8.74376  -6.223  0.00157 **
heights       0.71803     0.05154  13.932 3.42e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.521 on 5 degrees of freedom
Multiple R-squared:  0.9749,    Adjusted R-squared:  0.9699 
F-statistic: 194.1 on 1 and 5 DF,  p-value: 3.424e-05

```

Linear Regression: Weight vs. Height



```
Call:
glm(formula = am ~ mpg, family = binomial, data = mtcars)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-6.6035	2.3514	-2.808	0.00498 **
mpg	0.3070	0.1148	2.673	0.00751 **


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 43.230 on 31 degrees of freedom
Residual deviance: 29.675 on 30 degrees of freedom
AIC: 33.675

Number of Fisher Scoring iterations: 5

Mazda RX4	Mazda RX4 Wag	Datsun 710	Hornet 4 Drive
0.46109512	0.46109512	0.59789839	0.49171990
Hornet Sportabout	Valiant	Duster 360	Merc 240D
0.29690087	0.25993307	0.09858705	0.70846924
Merc 230	Merc 280	Merc 280C	Merc 450SE
0.59789839	0.32991148	0.24260966	0.17246396
Merc 450SL	Merc 450SLC	Cadillac Fleetwood	Lincoln Continental
0.21552479	0.12601104	0.03197098	0.03197098
Chrysler Imperial	Fiat 128	Honda Civic	Toyota Corolla
0.11005178	0.96591395	0.93878132	0.97821971
Toyota Corona	Dodge Challenger	AMC Javelin	Camaro Z28
0.49939484	0.13650937	0.12601104	0.07446438
Pontiac Firebird	Fiat X1-9	Porsche 914-2	Lotus Europa
0.32991148	0.85549212	0.79886349	0.93878132
Ford Pantera L	Ferrari Dino	Maserati Bora	Volvo 142E
0.14773451	0.36468861	0.11940215	0.49171990



RESULT:

Thus, to implement linear and logistic regression using machine learning is completed successfully.